

# mnist\_mlp-Copy1

August 2, 2018

## 0.1 Convolutional Neural Networks

---

In this notebook, we train an MLP to classify images from the MNIST database.

### 0.1.1 1. Load MNIST Database

```
In [1]: from keras.datasets import mnist

        # use Keras to import pre-shuffled MNIST database
        (X_train, y_train), (X_test, y_test) = mnist.load_data()

        print("The MNIST database has a training set of %d examples." % len(X_train))
        print("The MNIST database has a test set of %d examples." % len(X_test))
```

Using TensorFlow backend.

The MNIST database has a training set of 60000 examples.


The MNIST database has a test set of 10000 examples.

### 0.1.2 2. Visualize the First Six Training Images

```
In [2]: import matplotlib.pyplot as plt
        %matplotlib inline
        import matplotlib.cm as cm
        import numpy as np

        # plot first six training images

        fig = plt.figure(figsize=(20,20))
        for i in range(6):
            ax = fig.add_subplot(1, 6, i+1, xticks=[], yticks=[])
            ax.imshow(X_train[i], cmap='gray')
            ax.set_title(str(y_train[i]))
```




### 0.1.3 3. View an Image in More Detail

```
In [3]: def visualize_input(img, ax):
        ax.imshow(img, cmap='gray')
        width, height = img.shape
        thresh = img.max()/2.5
        print(thresh)
        for x in range(width):
            for y in range(height):
                ax.annotate(str(round(img[x][y],2)), xy=(y,x),
                            horizontalalignment='center',
                            verticalalignment='center',
                            color='white' if img[x][y]<thresh else 'black')

fig = plt.figure(figsize = (12,12))
ax = fig.add_subplot(111)
visualize_input(X_train[0], ax)
```

102.0



#### 0.1.4 4. Rescale the Images by Dividing Every Pixel in Every Image by 255

```
In [4]: # rescale [0,255] --> [0,1]
X_train = X_train.astype('float32')/255
X_test = X_test.astype('float32')/255
```

#### 0.1.5 5. Encode Categorical Integer Labels Using a One-Hot Scheme

```
In [5]: from keras.utils import np_utils

# print first ten (integer-valued) training labels
print('Integer-valued labels:')
print(y_train[:10])
```

```

# one-hot encode the labels
y_train = np_utils.to_categorical(y_train, 10)
y_test = np_utils.to_categorical(y_test, 10)

# print first ten (one-hot) training labels
print('One-hot labels:')
print(y_train[:10])

```

Integer-valued labels:

```
[5 0 4 1 9 2 1 3 1 4]
```

One-hot labels:

```

[[0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]

```

## 0.1.6 6. Define the Model Architecture

```

In [6]: from keras.models import Sequential
        from keras.layers import Dense, Dropout, Flatten

```

```

# define the model
model = Sequential()
model.add(Flatten(input_shape=X_train.shape[1:]))
model.add(Dense(512, activation='relu')) #512
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.2))

model.add(Dense(10, activation='softmax'))

# summarize the model
model.summary()

```

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_1 (Dense)	(None, 512)	401920

dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 10)	5130

=====

Total params: 669,706  
Trainable params: 669,706  
Non-trainable params: 0

-----

### 0.1.7 7. Compile the Model

```
In [7]: # compile the model
        model.compile(loss='categorical_crossentropy', optimizer='rmsprop',
                      metrics=['accuracy'])
```

### 0.1.8 8. Calculate the Classification Accuracy on the Test Set (Before Training)

```
In [8]: # evaluate test accuracy
        score = model.evaluate(X_test, y_test, verbose=0)
        accuracy = 100*score[1]

        # print test accuracy
        print('Test accuracy: %.4f%%' % accuracy)
```

Test accuracy: 11.9600%

### 0.1.9 9. Train the Model

```
In [9]: from keras.callbacks import ModelCheckpoint

        # train the model
        checkpointer = ModelCheckpoint(filepath='mnist.model.best.hdf5',
                                       verbose=1, save_best_only=True)
        hist = model.fit(X_train, y_train, batch_size=250, epochs=10,
                        validation_split=0.2, callbacks=[checkerpoint],
                        verbose=1, shuffle=True)
```

Train on 48000 samples, validate on 12000 samples

Epoch 1/10

48000/48000 [=====] - 6s 116us/step - loss: 0.3266 - acc: 0.8993 - va

Epoch 00001: val\_loss improved from inf to 0.12071, saving model to mnist.model.best.hdf5

```

Epoch 2/10
48000/48000 [=====] - 6s 116us/step - loss: 0.1266 - acc: 0.9608 - va

Epoch 00002: val_loss improved from 0.12071 to 0.09419, saving model to mnist.model.best.hdf5
Epoch 3/10
48000/48000 [=====] - 6s 121us/step - loss: 0.0846 - acc: 0.9732 - va

Epoch 00003: val_loss improved from 0.09419 to 0.08274, saving model to mnist.model.best.hdf5
Epoch 4/10
48000/48000 [=====] - 8s 158us/step - loss: 0.0649 - acc: 0.9797 - va

Epoch 00004: val_loss improved from 0.08274 to 0.07943, saving model to mnist.model.best.hdf5
Epoch 5/10
48000/48000 [=====] - 8s 174us/step - loss: 0.0519 - acc: 0.9828 - va

Epoch 00005: val_loss did not improve from 0.07943
Epoch 6/10
48000/48000 [=====] - 8s 159us/step - loss: 0.0419 - acc: 0.9865 - va

Epoch 00006: val_loss did not improve from 0.07943
Epoch 7/10
48000/48000 [=====] - 5s 107us/step - loss: 0.0329 - acc: 0.9888 - va

Epoch 00007: val_loss did not improve from 0.07943
Epoch 8/10
48000/48000 [=====] - 5s 104us/step - loss: 0.0293 - acc: 0.9901 - va

Epoch 00008: val_loss did not improve from 0.07943
Epoch 9/10
48000/48000 [=====] - 6s 116us/step - loss: 0.0277 - acc: 0.9907 - va

Epoch 00009: val_loss did not improve from 0.07943
Epoch 10/10
48000/48000 [=====] - 7s 136us/step - loss: 0.0232 - acc: 0.9922 - va

Epoch 00010: val_loss did not improve from 0.07943

```

### 0.1.10 10. Load the Model with the Best Classification Accuracy on the Validation Set

```

In [10]: # load the weights that yielded the best validation accuracy
         model.load_weights('mnist.model.best.hdf5')

```

### 0.1.11 11. Calculate the Classification Accuracy on the Test Set

```

In [11]: # evaluate test accuracy
         score = model.evaluate(X_test, y_test, verbose=0)
         accuracy = 100*score[1]

```

```

# print test accuracy
print('Test accuracy: %.4f%%' % accuracy)

```

Test accuracy: 98.0300%

### 0.1.12 Grid Search

```

In [ ]: import numpy
        from sklearn.model_selection import GridSearchCV
        from keras.models import Sequential
        from keras.layers import Dense
        from keras.wrappers.scikit_learn import KerasClassifier
        # Function to create model, required for KerasClassifier
        def create_model():
            # create model
            model = Sequential()
            model.add(Dense(12, input_dim=8, activation='relu'))
            model.add(Dense(1, activation='sigmoid'))
            # Compile model
            model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
            return model

        # fix random seed for reproducibility
        seed = 7
        numpy.random.seed(seed)

        # load dataset
        dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
        # split into input (X) and output (Y) variables
        X = dataset[:,0:8]
        Y = dataset[:,8]

        # create model
        model = KerasClassifier(build_fn=create_model, verbose=0)

        # define the grid search parameters
        batch_size = [10, 20, 40, 60, 80, 100]
        epochs = [10, 50, 100]
        param_grid = dict(batch_size=batch_size, epochs=epochs)
        grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
        grid_result = grid.fit(X, Y)

        # summarize results
        print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
        means = grid_result.cv_results_['mean_test_score']
        stds = grid_result.cv_results_['std_test_score']
        params = grid_result.cv_results_['params']
        for mean, stdev, param in zip(means, stds, params):
            print("%f (%f) with: %r" % (mean, stdev, param))

```